

Cache Misses

- On cache hit, CPU proceeds normally
- On cache miss
 - Stall the CPU pipeline
 - Fetch block from next level of hierarchy
- Instruction cache miss
 - Restart instruction fetch
- Data cache miss
 - Complete data access

CSULB

Measuring Cache Performance

- Components of CPU time
 - Program execution cycles
 - Includes cache hit time
 - Memory stall cycles
 - Mainly from cache misses
- With simplifying assumptions:

Memory stall cycles

$$= \frac{\text{Memory accesses}}{\text{Program}} \times \text{Miss rate} \times \text{Miss penalty}$$

$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Misses}}{\text{Instruction}} \times \text{Miss penalty}$$

CSULB

Cache Performance Example

- Given
 - I-cache miss rate = 2%, D-cache miss rate = 4%
 - Miss penalty = 100 cycles, Base CPI (ideal cache) = 2
 - Load & stores are 36% of instructions
 - Actual CPI = base CPI + I-cache cycles + D-cache cycles
- Miss cycles per instruction
 - I-cache: $0.02 \times 100 = 2$
 - D-cache: $0.36 \times 0.04 \times 100 = 1.44$
- Actual CPI = $2 + 2 + 1.44$
- Ideal CPU is $5.44 / 2$ times faster

CSULB

Average Access Time

- Average memory access time (AMAT)
 - **AMAT = Hit time + Miss rate \times Miss penalty**
 - At some point, the increase in hit time for a larger cache will overcome the improvement in hit rate leading to a decrease in performance
- Example
 - CPU with 1ns clock, hit time = 1 cycle, miss penalty = 20 cycles, I-cache miss rate = 5%
 - AMAT = $1 + 0.05 \times 20 = 2$


CSULB

Performance Summary

- When CPU performance increased
 - Miss penalty becomes more significant
- Decreasing base CPI
 - Greater proportion of time spent on memory stalls
- Increasing clock rate
 - Memory stalls account for more CPU cycles
- Can't neglect cache behavior when evaluating system performance

CSULB

Hotel California



Agustin
Avitia
Buendia
Coleman


Cuellar
Denholm
Espinoza
Han
Huang

Minaya
Moreno
Nguyen
Pham

Phung
Tecson
Tran
Vu
Yip

CSULB

Hotel California




| | |
|---------|----------|
| Agustin | Cuellar |
| Avitia | Denholm |
| Buendia | Espinoza |
| Coleman | Han |
| | Huang |

| | |
|--------|--------|
| Minaya | Phung |
| Moreno | Tecson |
| Nguyen | Tran |
| Pham | Vu |
| | Yip |

CSULB

Associative Cache Example



CSULB

Associative Caches

- Fully associative
 - Allow a given block to go in any cache entry
 - Requires all entries to be searched at once
 - Comparator per entry (expensive)
- n -way set associative
 - Each set contains n entries
 - Block number determines which set
 - (Block number) modulo (#Sets in cache)
 - Search all entries in a given set at once
 - n comparators (less expensive)

CSULB

Spectrum of Associativity

- For a cache with 8 entries

One-way set associative
(direct mapped)

| Block | Tag | Data |
|-------|-----|------|
| 0 | | |
| 1 | | |
| 2 | | |
| 3 | | |
| 4 | | |
| 5 | | |
| 6 | | |
| 7 | | |

Two-way set associative

| Set | Tag | Data | Tag | Data |
|-----|-----|------|-----|------|
| 0 | | | | |
| 1 | | | | |
| 2 | | | | |
| 3 | | | | |

Four-way set associative

| Set | Tag | Data | Tag | Data | Tag | Data | Tag | Data |
|-----|-----|------|-----|------|-----|------|-----|------|
| 0 | | | | | | | | |
| 1 | | | | | | | | |

Eight-way set associative (fully associative)

| Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data |
|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|
| | | | | | | | | | | | | | | | |

CSULB

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
- Direct mapped

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|---|--------|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | 0 | M | Mem[0] | | | |
| 8 | 0 | M | Mem[8] | | | |
| 0 | 0 | M | Mem[0] | | | |
| 6 | 2 | M | Mem[0] | | Mem[6] | |
| 8 | 0 | M | Mem[8] | | Mem[6] | |

CSULB

Associativity Example

- 2-way set associative

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|--------|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | 0 | M | Mem[0] | | | |
| 8 | 0 | M | Mem[0] | Mem[8] | | |
| 0 | 0 | H | Mem[0] | Mem[8] | | |
| 6 | 0 | M | Mem[6] | Mem[8] | | |
| 8 | 0 | H | Mem[6] | Mem[8] | | |

CSULB

Associativity Example

▪ Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|--------|--------|--|
| 0 | | M | Mem[0] | | | |
| 8 | | M | Mem[0] | Mem[8] | | |
| 0 | | H | " | " | | |
| 6 | | M | " | " | Mem[6] | |
| 8 | | H | " | " | " | |

CSULB

How Much Associativity

- Increased associativity decreases miss rate
 - But with diminishing returns
- Simulation of a system with 64KB D-cache, 16-word blocks, SPEC2000
 - 1-way: 10.3%
 - 2-way: 8.6%
 - 4-way: 8.3%
 - 8-way: 8.1%

CSULB

CSULB



CSULB

- CSULB

Multilevel Caches

- Primary cache attached to CPU
 - Small, but fast
- Level-2 cache services misses from primary cache
 - Larger, slower, but still faster than main memory
- Main memory services L-2 cache misses
- Some high-end systems include L-3 cache

CSULB

Multilevel Cache Example

- Given
 - CPU base CPI = 1, clock rate = 4GHz
 - Miss rate/instruction = 2%
 - Main memory access time = 100ns
- With just primary cache
 - Miss penalty = $\frac{100}{0.25} = 400$
 - Effective CPI = $1 + 0.02 \times 400 = 9$

CSULB

Example (cont.)

- Now add L-2 cache
 - Access time = 5ns
 - Global miss rate to main memory = 0.5%
- Primary miss with L-2 hit
 - Penalty = $\frac{5}{0.25} = 20$
- Primary miss with L-2 miss
- CPI = $1 + 0.02 \times 20 + 0.005 \times 400 = 3.4$
- Performance ratio = $9/3.4$

CSULB

Multilevel Cache Considerations

- Primary cache
 - Focus on minimal hit time
- L-2 cache
 - Focus on low miss rate to avoid main memory access
 - Hit time has less overall impact
- Results
 - L-1 cache usually smaller than a single cache
 - L-1 block size smaller than L-2 block size

CSULB

Virtual Machines

- Host computer emulates guest operating system and machine resources
 - Improved isolation of multiple guests
 - Avoids security and reliability problems
 - Aids sharing of resources
- Virtualization has some performance impact
 - Feasible with modern high-performance computers
- Examples
 - IBM VM/370 (1970s technology!)
 - VMWare
 - Microsoft Virtual PC

CSULB

Virtual Machine Monitor

- Maps virtual resources to physical resources
 - Memory, I/O devices, CPUs
- Guest code runs on native machine in user mode
 - Traps to VMM on privileged instructions and access to protected resources
- Guest OS may be different from host OS
- VMM handles real I/O devices
 - Emulates generic virtual I/O devices for guest

CSULB

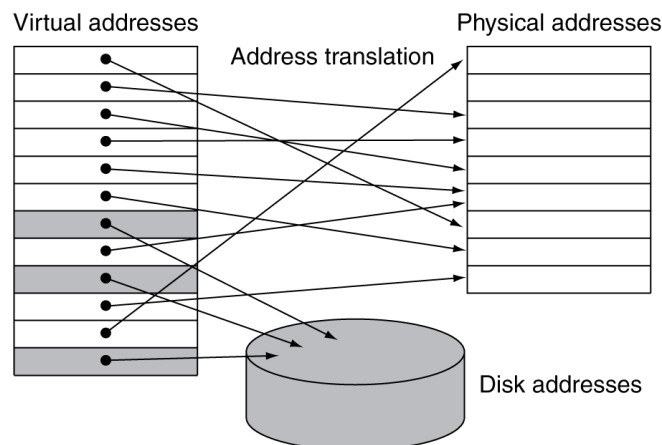
Virtual Memory

- Use main memory as a “cache” for secondary (disk) storage
 - Managed jointly by CPU hardware and the operating system (OS)
- Programs share main memory
 - Each gets a private virtual address space holding its frequently used code and data
 - Protected from other programs
- CPU and OS translate virtual addresses to physical addresses
 - VM “block” is called a page
 - VM translation “miss” is called a page fault

CSULB

Address Translation

- Fixed-size pages (e.g., 4K)



CSULB

Page Fault Penalty

- On page fault, the page must be fetched from disk
 - Takes millions of clock cycles
 - Handled by OS code
- Try to minimize page fault rate
 - Fully associative placement
 - Smart replacement algorithms